

Department of Computer Science

University of Kaiserslautern

Master Thesis

Offline caching in web applications for AntidoteDB

Server Khalilov

contact@serverkhalilov.com

University of Kaiserslautern
Department of Computer Science
Software Engineering

Leader:
Prof. Dr. Arnd Poetzsch-Heffter

Supervisor:
Dr. rer. nat. Annette Bieniusa



Zusammenfassung

Zusammenfassung auf deutsch

Abstract

The purpose of this thesis is to explore the possibilities of developing an offline web application, which would serve as a client for the AntidoteDB database. We developed a prototype of the application and designed the architecture of the application in such a way that both offline and online functionalities are possible.

To model the data stored offline in the local database of a web-browser, Conflict-free Replicated Data Types (CRDTs). It lets to ease the task of merging the changes.

Apart from that, the client-server protocol was designed, in order to support the functionality of the application. The paper could be divided into two parts: firstly, the problem of designing mentioned solution is going to be discussed. Secondly, there is an implementation part and a description of how specific problems were tackled.

Acknowledgement

I would like to take this opportunity and express my biggest gratitude to my friends and everyone, who supported me through all the ups and downs I had during my time at the University of Kaiserslautern.

Firstly, I am very grateful to Prof. Dr. Arnd Poetzsch-Heffter and my thesis supervisor Dr. rer. nat. Annette Bieniusa for making it possible for me to work on the topic that suits my area of interest. Moreover, I want to thank Dr. rer. nat. Annette Bieniusa for having the door to her office always open for a discussion, whenever I needed it. I could not ask for more.

Next, I would like to thank the staff at Software Technology Group, who made me feel very welcome and patiently answered any questions I had. Especially, Mathias Weber, Peter Zeller, and Deepthi Akkoorath.

Finally, I want to mention my family and, particularly, my parents. They did their very best to encourage and motivate me throughout the whole period of my study in Germany. Nothing of it would ever be possible without them. Thank you.

Ich versichere hiermit, dass ich die vorliegende Masterarbeit mit dem Thema „Offline caching in web applications for AntidoteDB“ selbstständig verfasst und keine anderen als die angegebenen Hilfsmittel benutzt habe.

Die Stellen, die anderen Werken dem Wortlaut oder dem Sinn nach entnommen wurden, habe ich durch die Angabe der Quelle kenntlich gemacht.

Kaiserslautern, den 15. December 2018

Server Khalilov

Contents

1	Introduction	1
1.1	Motivation	1
1.2	Research questions	1
1.3	The structure	2
2	Theoretical background	3
2.1	Main concepts	3
2.2	AntidoteDB	5
2.3	Conflict-Free Replicated Datatypes (CRDTs)	6
3	Design	11
3.1	Requirements	11
3.2	Assumptions	11
3.3	Protocol	13
3.3.1	Data transmission	13
3.3.2	Description	13
3.3.3	A client sends an update to the server	16
3.3.4	Two clients interact with a server.	17
3.4	Offline functionality	18
3.5	Online functionality	19
3.6	The transition between offline and online modes	19
4	Technologies	21
4.1	Service Worker	21
4.1.1	Service worker lifecycle	23
4.2	Cache API	25
4.3	IndexedDB	26
4.3.1	IndexedDB terms	27
4.3.2	IndexedDB Promised	27
4.4	Persistent Storage	28
4.5	Background Sync	28

Contents

5	Implementation	31
5.1	System's main components	32
5.1.1	Web Application	32
5.1.2	Server	33
5.1.3	A database layer	33
5.2	Optimization	34
6	Evaluation	35
7	Related Work	37
7.1	SwiftCloud: a transactional system that brings geo-replication to the client	37
7.2	Legion: a framework, which enriches web applications	37
7.3	Developing web and mobile applications with offline usage experience	38
8	Conclusion	45
8.1	Summary	45
8.2	Future Work	45
	List of Figures	47
	List of Tables	49
	List of Code	51
	Listings	53
	Bibliography	55

1 Introduction

In this chapter we are going to discuss the motivation, research questions and the scope of this thesis.

1.1 Motivation

The main motivation of this thesis is to explore the possibilities of implementing a web-client for the AntidoteDB with a support of caching and by utilizing the main features of the database. As AntidoteDB is already known for its use in the development of applications, it would dramatically improve the user-experience if offline work with the database is going to be reached. We are going to explore the best ways to develop such an application, design it and implement. The goal of this thesis is to show you that it is possible to build an application based on Antidote-DB that works both online and offline. Here and from now on, we will call the application we are going to develop – WebCure.

1.2 Research questions

The following research questions are going to be addressed in this thesis:

- **RQ1.** How efficient is it to use a web-client with cache rather than without it?
- **RQ2.** What are the methods available to implement web-applications that would be able to work off-line and in the conditions of poor network connections?
- **RQ3.** What could be a scalable solution for transmitting CRDT data between a server and clients?

1.3 The structure

Here, I am going to briefly sketch the structure of this thesis and explain the contents of each chapter.

- Description of the main requirements of the application;
- Design of the architecture
- Description of the implementation phase
- Evaluation
- Conclusion

2 Theoretical background

In this chapter, we are going to introduce the reader to the theoretical concepts, which represent a prerequisite to have the understanding of this thesis.

2.1 Main concepts

Distributed database is “a collection of multiple, logically interrelated databases distributed over a computer network”[1]. A *geo-distributed database*, in its turn, is a database, which is spread across two or more geographically distinct locations and runs without experiencing performance delays. The maintaining of such databases brings its challenges. As the database is spread across several locations, there should be a replication process in order to ensure that replicas of that database synchronize and have the latest state of the data. This replication process should be fast, because if there are two replicas of the database, then whenever there is some information written to the first replica, it should be accessible to users, who use the second one. That is the problem of the availability, but before the information at replicas becomes available, it first should be checked over the consistency, as the states of the replicas should be equal. That is a complex task to solve.

Working with such a distributed database, whenever the data is needed to be read or changed in any way, a transaction should be started, executed and closed. A *transaction* is a basic unit of computing, which consists of sequence of operations that are applied atomically on a database. Transactions transform a consistent database state to another consistent database state. A transaction is considered to be *correct*, if it obeys the rules, specified on the database. As long as each transaction is correct, a database guarantees that concurrent execution of user transactions will not violate database consistency [1]. *Consistency* requires transactions to change the data only according to the specified rules. An example of consistency rule can be the following: let us say that in a bank database the bank account number should only consist of integer numbers. If an employee tries to create an account that contains something other than

2 Theoretical background

integer numbers in it, then the database consistency rule will disallow it. Consistency rules are important as they control the incoming data and reject the information, which does not fit.

Sequential consistency and *linearizability* are two consistency conditions that are well-known. Sequential consistency requires that all of the data operations appear to have executed atomically (i.e. independently), in some sequential order. When this order must also preserve the global ordering of non-overlapping operations, this consistency is called linearizability[2]. Linearizability guarantees that the same operations are applied in the same order to every replica of the data item[3]. Serializability is a guarantee about transactions, that they are executed serially (i.e. without overlapping in time) on every set of the data items[3]. Serializability is more strict than sequential consistency, as the granularity of sequential consistency is a single operation, while for serializability it is a transaction. As a result, when serializability is satisfied, the sequential consistency is also satisfied, but not vice versa.

Now, let us introduce different consistency models and the one we will follow in the designing part of WebCure.

As stated by Shapiro [4], “*strong consistency* model could be described in the following way: whenever the update is performed, everyone knows about it”. It means that there is a total order of updates and reads are guaranteed to return the latest data, irregardless of which replica is the source of the response. The advantage of strong consistency is that the database is always in a consistent state and to disadvantages we can add low latency, as there is a delay for making sure that all the replicas are in a consistent state before any other read / write requests could be processed. The latency point is a huge drawback for the performance, especially if strong consistency model is considered to be used as a solution for the web, where users usually expect high responsiveness and availability.

The main point of replicating data is to improve such aspects as reliability, availability, performance and latency. However, according to CAP Theorem[5], a distributed database can only have two of the three properties: consistency, availability and partition tolerance. This theorem is very important, as it makes people think towards the trade-off between those three properties for a specific use case. There are some weaker consistency models, where the results of requests can alter depending on the replica[6]. In this thesis, we will stick with partial *causal consistency*. As it is stated in Zawirski et al. [7], causal consistency is “the strongest available and convergent model”. They continue their statement saying that “under causal consistency, every process observes a monotonically non-decreasing set of updates that includes its own updates, in an order that

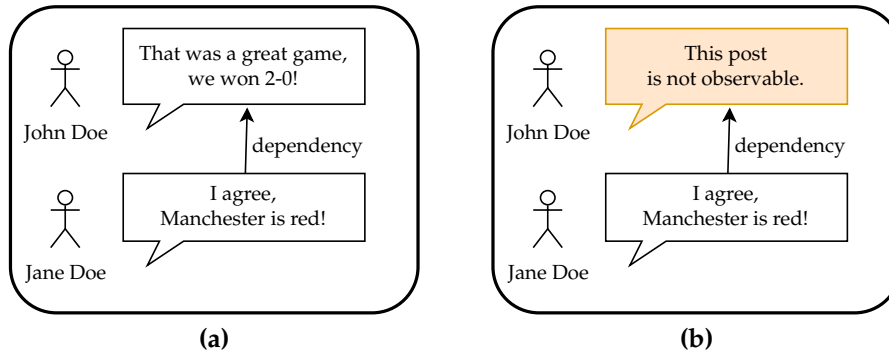


Figure 2.1: An example of how a causally consistent behaviour (a) and the one that is not (b) could work in a social network.

respects the causality between operations". As the causal ordering is respected, it makes it easier for programmers to reason, as it gives the guarantee that related events are visible in the order of occurrence, while the events, which have no relation to each other, can be in different order in different replicas. Let us consider an example of an application for some of social networks. There, a reply to a wall post happens after the original post is published. Thus, users should not see the reply before the original post is observable. This type of guarantees are provided by causal consistency. Looking at the [Figure 2.1](#) you can see that on the left subfigure, the user can see the original wall post as well as the reply, while on the right subfigure, without causal consistency, the user sees only the reply, while the original wall post is missing.

2.2 AntidoteDB

For this thesis, one of the core parts in the architecture of the WebCure belongs to the database called AntidoteDB[8]. It helps programmers to write correct applications and has the same performance and horizontal scalability as AP / NoSQL[9], while it also:

- is geo-distributed, which means that the datacenters of AntidoteDB could be spread across anywhere in the world;
- groups operations into atomic transactions[10, 11];
- delivers updates in a causal order and merges concurrent operations;

2 Theoretical background

Merging concurrent operations is possible because of Conflict-Free Replicated Datatypes (CRDTs) [12], which are used in AntidoteDB. It supports counters, sets, maps, multi-value registers and other types of data that are designed to work correctly in the presence of concurrent updates and failures. The usage of CRDTs allows the programmer to avoid problems that are common for other databases like NoSQL, which are fast and available, but hard to program against[11]. We will cover the topic of CRDTs later in this chapter.

Apart from that, to replicate the data AntidoteDB implements the *Curere*[11] protocol. It is a highly scalable protocol, which provides causal consistency.

To ensure the guarantees it offers, AntidoteDB uses timestamps, indicating the time after the transaction. Timestamps are considered to be unique, totally ordered, and consistent with causal order, which means that if *operation 1* happened before *operation 2*, then the timestamp related to the *operation 2* is greater than the one related to the *operation 1*[12]. Whenever the update operation has to be applied, it is also possible to provide a minimum time from what that update should be performed. This information is useful, when a client is working with a server, which is based on AntidoteDB. In such cases, when one data center stops working, the client can reconnect to another one. As a client can remember the latest timestamp for the data it has worked on, the failover to another data center is possible without any additional efforts, as the timestamp information will help the client to request only the data, which has timestamps greater than the one, which is already stored at the client.

2.3 Conflict-Free Replicated Datatypes (CRDTs)

As it is stated in the work of Preguiça et al. [13], a conflict-free replicated datatype (CRDT) is an abstract datatype, which is designed for a possibility to be replicated at multiple processes and possesses the following properties:

- The data at any replica can be modified independently of other replicas
- Replicas deterministically converge to the same state when they received the same updates

Replication is a fundamental concept of distributed systems, well studied by the distributed algorithms community[12]. There are two models

of replication that are considered: state-based and operation-based. We are going to introduce our reader to both of them below.

Operation-based replication approach




Figure 2.2: Operation-based approach[12]. «S» stands for source replicas and «D» for downstream replicas.

In this thesis, we are going to use the operation-based replication approach, where replicas converge by propagating operations to every other replica[13]. Once an operation is received in a replica, it is applied locally. Afterwards, all replicas would possess all of the updates. At the **Figure 2.2**, you can see that firstly operations $f(x_i)$, $g(x_i)$ applied locally at source replicas x_i , and then the operations are conveyed to all the other replicas. The second part of this process is named *downstream* execution.

This replication approach infers that replicas do not exchange full states with each other, which is a positive in terms of efficiency. Not always, though, as it depends on the task. Sometimes, applying multiple operations at every replica could be costly and this is where state-based replication approach is beneficial.

State-based replication approach

The idea of this approach is kind of opposite to the operation-based one. Here, every replica, when it receives an update, first applies it locally. Afterwards, it sends its updated state to other replicas. Following this way, every replica sends its current full state to other replicas. Afterwards, the merge function is applied between a local state and a received state, and

2 Theoretical background



Figure 2.3: State-based approach[12]. «S» stands for source replicas and «M» for merging stages.

every update eventually is going to appear at every other replica in the system. You can see this at the **Figure 2.3**, where x_i stands for replicas, $f(x_i)$, $g(x_i)$ stands for the functions that apply updates locally at source replicas before sending a new full state for a follow-up merges to other replicas in the system.

There are different types of CRDTs, however, we will consider three of them: counters, sets and multi-value registers. We will give a brief description to each of the mentioned datatypes below.

Counter

This is a datatype, which keeps track on its state, which is an integer number. The value of the Counter could be modified by the operations *inc* and *dec* that increases or decreases the state by one unit, accordingly[13]. The concurrency semantics for this datatype is that a final state of the object reflects all the performed operations on it. In other words, to calculate the state of the Counter, it is needed to count the number of increments and subtract the number of decrements.

Add-wins Set

Add-wins Set¹ datatype represents a collection of objects with a specific handling of concurrent updates performed over them. In case of concur-

¹for the simplicity of reading, later it goes as Set

2.3 Conflict-Free Replicated Datatypes (CRDTs)

rent updates on the same object, *add* operations in Set win against *remove* operations. If, for example, there is an empty set and two concurrent operations applied to it – *add a* and *remove a*, then the result is going to be $\{a\}$, as *add* operation wins. A *remove* operation will “overwrite” an *add* operation only when it happens after it[13].

Multi-value Register

This datatype maintains a value and provides a *write* operation of updating that value. The interesting part about Multi-value Registers is their concurrency semantics. In case of two or more updates happening at the same time, all values are kept. Thus, the state of the register will consist of all the concurrently written updates for a further processing. However, any additional single *write* operation will overwrite the previous state of the register, even if it consisted of multiple values[14].

3 Design

In this chapter, we are going to, firstly, introduce the requirements and assumptions we make for the WebCure. Having set them, afterwards, we will in details describe the design of the solution.

3.1 Requirements

We listed the functional requirements of the WebCure at **Table 3.1** and non-functional requirements at **Table 3.2**.

3.2 Assumptions

In this section, we are going to give a list of assumptions, which makes the WebCure work correctly.

1. **Timestamps.** Firstly, the database storage used for the server's side should have the concept of timestamps (like in AntidoteDB, described in **Section 2.2**), in order for the protocol we are going to describe in the **Section 3.3** to work correctly.
2. **Cache is persistent.** For the WebCure to work properly online and, especially, offline, we assume that the browser's cache is safe from automatic clearing. Otherwise, if the cache could be cleared automatically depending on the browser's behaviour, it makes it impossible to support the claim that the application can work offline. To guarantee this assumption, a Persistent Storage API described in **Section 4.4** can be used.
3. **Name duplicates.** We restrict the creation of different CRDT elements with the same name in the system due to limitations of AntidoteDB, as it is described by the requirement *R7* in the **Table 3.1**. As AntidoteDB is in the process of ongoing development, currently the database crashes when there is an attempt to create elements of

Table 3.1: Functional requirements.

R1	Retrieval, increment and decrement of the Counter CRDT should be possible.
R2	Retrieval, adding and removing elements from the Set CRDT should be possible.
R3	Retrieval, assigning and resetting the Multi-Value Register CRDT should be possible.
R4	Retrieval elements of any supported CRDTs should be possible according to the passed timestamp.
R5	Do not store the element's values in the cache, if they were requested at specific timestamp (for the matter of having only latest data on the client's side).
R6	The user should be able to remove from the client any stored data element.
R7	It should not be possible to create elements of different CRDTs with the same name (due to limitations of AntidoteDB)
R8	When offline, it should be possible to perform operations on supported CRDTs.
R9	Any operations performed offline, once the connection is restored, should be sent to the server immediately.
R10	When the connection is re-established after having data changes in offline mode, the client data should be updated with relation to current client's changes, as well as with relation to the changes made on the server.

Table 3.2: Non-functional requirements.

NFR1	The web application should be available online and offline (except of the functionality with timestamp-related updates).
NFR2	The web application should be available with a poor internet connection.

different CRDTs with the same name. Thus, this condition has to be satisfied.

4. **Server's database is always on.** We assume that the server's database is not going to be reset and lose its all data. As the client completely relies on the server's storage for the synchronization and only sends back operations performed offline on client's side, it will be impossible to restore the server's database from the client's storage, even if it was up to date before the server's data loss. With additional changes to the current protocol it might be possible, though, but that is not the topic we cover in this thesis. However, even in such situation, the client will be able to continue the offline work.

As we specified the requirements, we can go further into and design the protocol of the system.

3.3 Protocol

The fundamental part of the WebCure will be its protocol design. We are going to describe it in an event-based way in the form of pseudo-code in the following sections.

3.3.1 Data transmission

As we already know from the **Chapter 2**, because the AntidoteDB is using CRDT datatypes, the following options are possible to update the database: state-based and operation-based. Due to the time constraints and the amount of work, this thesis will consider only the operation-based approach. Therefore, whenever a client wants to update the database, it will send to the server a list of operations. However, whenever it wants to read the value, it will receive the current state of the object from the database. For this thesis, we are going to use such datatypes as Counters, Sets and Multi-Value Registers, to which the reader was introduced in the **Section 2.3**.

3.3.2 Description

Firstly, as we would like to concentrate on the communication part between a server and a client, let us for now keep both of them as block

3 Design




Figure 3.1: An overview of the communication protocol.

boxes², as they are represented at **Figure 3.1**. Next, we will go through different stages of their communication and describe, how we treated these processes.

Graphics notations




Figure 3.2: An overview of notations used in the following chapters for the protocol explanation.

Let us explain the notations, which are going to be used for a further

²In software engineering, a black box is a system, which can be viewed in terms of its inputs and outputs, without the understanding of its internal workings.[15]

protocol description. At **Figure 3.2 (a)**, we can see a notation for the timeline, every point of which represents a different timestamp. Timelines will be used for the matter of showing the sequence of events happening. At **Figure 3.2 (b)**, the arrow shows the transmission of data between a system *A* and system *B*, as well as its direction and a command. **Figure 3.2 (c)** represents the state of the data at a specific point in time, while **Figure 3.2 (d)** points to a timestamp, at which an operation that changes the system's storage was applied.

Next, as we already mentioned, we will explain the protocol in an event-based way.

A client receives an update from the server




Figure 3.3: The way a client receives a key-value update from the server.

Let us assume that a client starts its work with an empty storage. Then, a user might want to request the actual data from the server. In this case, as can be seen at **Figure 3.3**, a user has to pass to the server an *id* of the data to request. If the request is successful, the server is going to respond with a *value* for the requested *id* and the timestamp of the last write – t_0 , so the client will store this information on its own storage.

Listing 3.1 A pseudocode for requesting the data on a client's side.

```

1 // Read function that pulls database changes
2 // @param id: the id of the object, for which the update was
   requested;

4 function read(id) {
5   if (ONLINE) {
6     value = fetch(id); // request the value of the data by id
                           from the server
  
```

3 Design

```
7     store(value, timestamp); // store the value and timestamp
    received from the server on a client's storage
8     show(value, timestamp); // show the value and timestamp
    received from the server to the user;
9 } else {
10    value = fetch(id); // Get the value of the data by id
    from the cache
11    show(value); // show the value to the user
12 }
13 }
```

The pseudocode of the logic for this *read* functionality can be seen at **Listing 3.1**. If a client is online, then the value and timestamp associated with passed *id* is going to be fetched from the server and, afterwards, stored on a client's storage and displayed to the user. However, if the client is offline, then the client will try to get the data from the cache for the requested *id*, or, if there is none, it will show the appropriate message to the user.

In case a client receives an update from the server, and that update has an earlier timestamp than the one, which is already stored in the client's cache, then a client skips it and tries later for fresher updates. The implementation of this will be explained later in the chapter of implementation.

3.3.3 A client sends an update to the server

In the case of writing the information to the server, a client has to send a key with an operation to the server. Then the acknowledgment with a timestamp is going to be sent back to the client.




Figure 3.4: The way a client sends a write operation to the server.

Several updates on a client while offline

The client is capable of performing updates when offline. These updates will take effect immediately. However, in order for them to be applied to the server, the client has to be back online. Once the connection is established again, all the updates that were performed on the client-side in offline mode will be sent to the server.

3.3.4 Two clients interact with a server.




Figure 3.5: Graphical representation over a possible communication between two clients and a server.

Let us assume that initially, the server has the key-value pair (k, v) at the timestamp t_0 . Therefore, after both clients receive updates from the server, they consist of the (k, v) pair at the timestamp t_0 . In case clients change the value under mentioned key to something else, they will have to get an acknowledgement from the server in order to receive a unique timestamp related to the change. At the representation above, a *Client 1* is acting first and getting an acknowledgement of its change at the time t_1 , while *Client 2* makes the change later at time t_2 . That makes *Client 1* receive a new value v'' , when it reads the information from the server again.

3.4 Offline functionality

In this section, we are going to describe the offline functionality of the system.

Initially, the database is empty. Therefore, if the user is offline from the very beginning, he should be able to add the data into the database himself. The system represented in **Figure 5.2** will change by having only the Web Application and the database. However, whenever the connection is established with the server, the operations, which were stored in local database while offline, will be sent to the server.

At the **Figure 3.6**, the sequence of getting the data from the local database is shown. This case describes the scenario, when the server is unavailable and the application has to read the value from the local database.

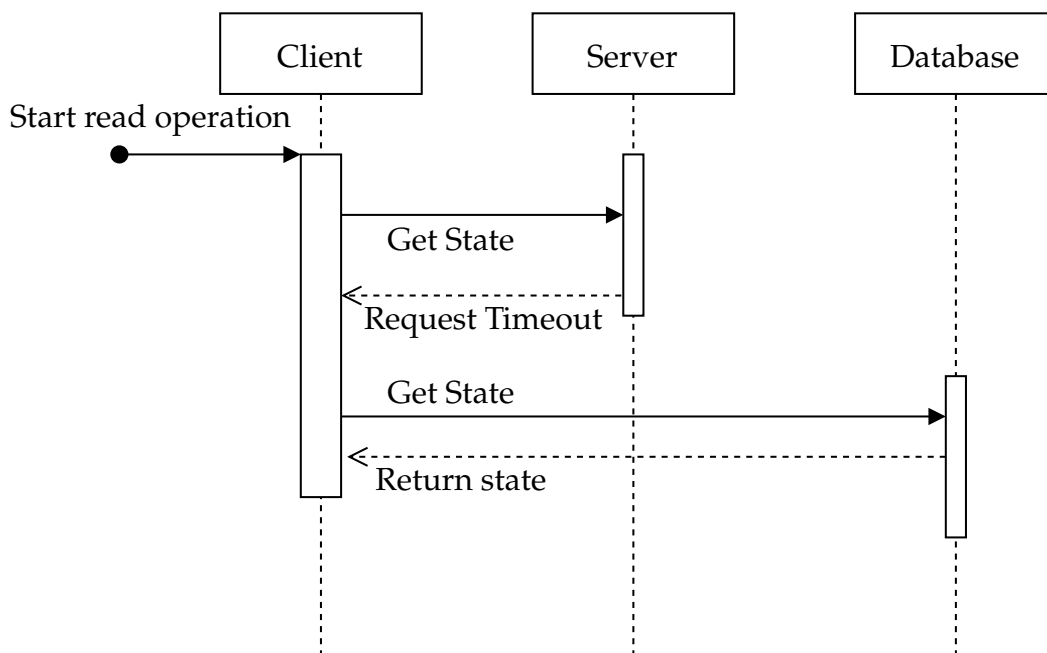


Figure 3.6: Successful request of state from the local database – Sequence diagram.

At the **Figure 3.7**, the sequence of storing the data locally is shown. When the connection is not there, the application will store in the database all offline-performed operations by the user. Afterwards, once the connection is re-established, that data will be sent back to the server. At the point when we read the state again from the server, the data, which is stored locally could be easily removed due to be pointless to hold on it.

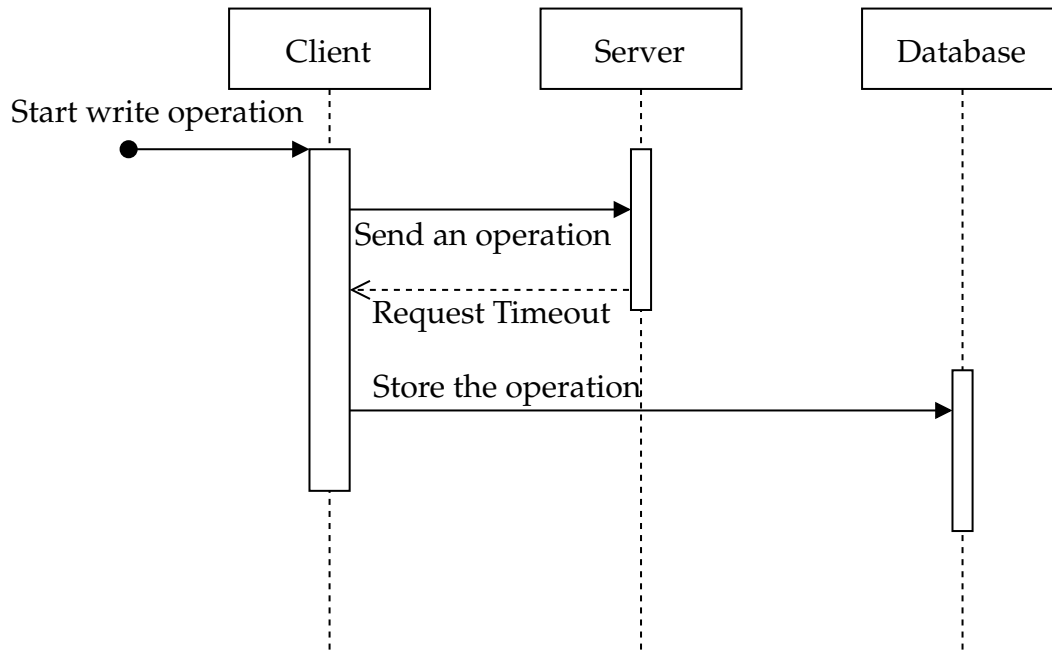


Figure 3.7: Successful storing of an operation in the local database – Sequence diagram.

3.5 Online functionality

In this section, we are going to describe the online functionality of the system.

3.6 The transition between offline and online modes

4 Technologies

In this chapter, we describe the technologies we used to implement the design we introduced in the **Chapter 3**. To give an overview, the implementation of WebCure’s design demands the following components:

- A service worker, in order to manage requests coming from the client;
- A Cache API to store HTML, CSS, JavaScript files, and any static files[16] to make the application available offline;
- A database to store the data locally;
- Background Sync for deferring the actions conducted offline until the connection is stable;

4.1 Service Worker

Service worker[17] is a web worker³, which is a JavaScript file, which lies in the middle, between the web application and network requests. Service worker runs independently from the web-page and has the following characteristics:

- It does not have access to the DOM⁴, however, it has the control over pages (not just a specific one);
- When the application is not running, a service worker still can receive push-notifications from the server[19], which let us improve the user experience of the application and makes it “closer” to native mobile applications;

³Web Workers make it possible to run a script operation in a background thread separate from the main execution thread of a web application[18].

⁴DOM, or **D**ocument **O**bject **M**odel of the page defines HTML elements as objects, as well as their properties, methods, and events.

4 Technologies

- It runs over HTTPS, as for the projects using a service worker, the “man-in-the-middle”⁵ attacks could represent a real threat;
- It offers a possibility to intercept requests as the browser makes them and has the following options:
 - to let requests go to the network as usual;
 - to skip the network and redirect requests to get the data from the cache;
 - to perform any combination of the above;




Figure 4.1: An overview of the service worker being able to provide the application experience of both online and offline modes.

We are going to use a service worker in order to intercept the network traffic. For example, as can be seen at **Figure 4.1**, in circumstances when problems are happening with a network when a request is in processing, we will get the data from the cache. At times, when the internet connection could not be established, the cache content can be easily displayed without forcing a user to wait, which dramatically improves performance and user experience. Also, when the internet connection is present, the data can be received from the actual network.

⁵In computer security, this is a type of attack, where a third party secretly alters the communication between two parties, while they believe they are directly “talking” to each other[20].

4.1.1 Service worker lifecycle

In this subsection, we will introduce the steps through which a service worker goes. These steps are registration, installation, and activation. All of them we are going to describe below in more details.

Registration

Before one can use service worker features, a developer has to register the corresponding script in the JavaScript code. The registration helps the browser to find a service worker and, afterwards, start its installation in the environment.

Listing 4.1 An example code, which demonstrates how to register a service worker[19].

```
1 if ('serviceWorker' in navigator) {  
2   navigator.serviceWorker.register('/service-worker.js')  
3   .then(function(registration) {  
4     console.log('Registration successful, scope is:',  
6       registration.scope);  
5   })  
6   .catch(function(error) {  
7     console.log('Service worker registration failed, error:',  
9       error);  
8   });  
9 }
```

Looking at the *line 1* of code at **Listing 4.1**, we see that, firstly, it is necessary to check, whether the browser supports service workers by observing the property *serviceWorker* of *navigator*⁶. If so, the service worker is going to be registered then, as can be seen at the *line 2*, where the location of the service worker is stated as well. The *navigator.serviceWorker.register* returns a promise, which resolves when the registration was successful. Afterward, the scope of the service worker is logged with *registration.scope*.

Not every browser supports the functionality of service workers. Nowadays, the latest versions of such browsers as Chrome, Edge (partially)

⁶The navigator object contains information about the browser[21].

and Firefox support it. However, it is still only partially supported in Opera and Safari, and not at all in the Internet Explorer[22] – that is something developers should keep in mind.

The scope of the service workers is quite significant: it defines the paths, from which it can intercept network requests. By default, the scope of the service worker is the location, where the service worker file is stored, including all the sub-directories. For example, if the scope is the root directory, then the service worker is going to regulate the requests for all files at the domain.

Listing 4.2 An example code, which demonstrates how to set a custom scope when registering a service worker[19].

```
1 navigator.serviceWorker.register('/service-worker.js', {  
2   scope: '/app/'  
3 });
```

When registering, it is also possible to use a custom scope. **Listing 4.2** demonstrates that the service worker is going to have a scope of `/app/`. It indicates that it will control requests from all the pages like `/app/` and deeper. When the service worker is already installed, `navigator.serviceWorker.register` returns the object of the currently active service worker.

Installation

After the registration, the browser might attempt installation of a service worker, if it is considered as a new one, which happens in the following situations:

- when the site does not have a registered service worker yet;
- when there is a difference between the previously installed service worker and the new one;

Installation triggers service worker's *install* event. There is a probability of having a listener for it in order to assign a specific task (depends on the use case), which follows the installation on success. The way to set up this listener is shown at **Listing 4.3**.

Listing 4.3 An example code, which demonstrates a listener for the *install* service worker's event[19].

```
1 // Listen for install event, set callback
2 self.addEventListener('install', function(event) {
3     // Perform some task
4 });
```

Activation

After the installation, a service worker has to be activated. In case there are open pages, which are controlled by the previous version of a service worker, then the recently installed service worker will be waiting. The activation of a new service worker takes place only when there are no other pages, controlled by the old version of a service worker. That provides a guarantee that only one version of service worker manages the pages of its scope at any time.

Similar to the installation, the activation phase also has its event that gets triggered once the service worker is activating, as we can see at **Listing 4.4**. For example, at this point, it might be a good idea to clean the previously stored old cached data.

Listing 4.4 An example code, which demonstrates a listener for the *activation* service worker's event[19].

```
1 self.addEventListener('activate', function(event) {
2     // Perform some task
3 });
```

4.2 Cache API

For the best offline experience, the web application should store somewhere HTML, CSS, JavaScript code, as well as images, fonts. There is a place for all of it called Cache API, which we are going to use for Web-

Cure. With Cache API it is possible to store network requests associated with corresponding requests.

Listing 4.5 An example code, which demonstrates how one can create cache storage called *my-cache*[23].

```
1 caches.open('my-cache').then((cache) => {  
2   // do something with cache...  
3 });
```

At **Listing 4.5** we see an example code of how a cache with a name *my-cache* is created. If the operation is performed successfully, then a promise⁷ resolves and one is going to get access either to a newly created cache or to the one, which existed before the call of *open* method.

That covers the main functionality of the API. After creating a cache, it is possible to manipulate it in many ways.

Let us now introduce one of the operations – adding elements to the cache. A developer can provide a string for the URL that will be fetched and stored in a cache object. Whenever the data is received from the cache, the browser will return a particular object according to the URL that was stored in the cache. More details on how long the data could be stored in the cache are given in **Section 4.4**.

Apart from adding to the cache object, it is also possible to remove stored URLs, check the current list of cached requests, delete a cache object, and other operations.

4.3 IndexedDB

As the protocol we introduced in the **Chapter 3** clearly describes that there is a necessity to have a client-side storage system, we are going to use for that purpose IndexedDB[25]. It is a large-scale, NoSQL database, which lets us store any data in the user's browser. Apart from that, it supports transactions and achieves a high-performance search due to the usage of indexes.

⁷JavaScript Promise is an object, which represents the eventual completion or failure of an asynchronous operation[24].

4.3.1 IndexedDB terms

In order to properly understand how IndexedDB works, it is quite useful to understand the concepts that are used in the database. First of all, each *IndexedDB database* contains *Object stores*. Those object stores, in its turn, are similar to tables in traditional databases. Usually, the practice is to have one object store for each type of data. This data could be anything: custom objects, strings, numbers, arrays, and other. It is possible to create more than one database, which could contain various object stores, but normally it is one database per application, which should have one object store for each type of data stored. To expedite the way of identifying objects in object stores, the latter have *primary keys*, which must be unique in the particular object stores. Primary keys are defined by the developer and are very useful regarding searching the data.

All read and write operations in IndexedDB should be wrapped into a *transaction*, which guarantees the database integrity. The critical point is that if one operation within a transaction fails, then none of the other operations are going to be applied.

4.3.2 IndexedDB Promised

As IndexedDB is relatively a new API, it is not supported by all the web browsers yet and, therefore, the support for it should be checked before any further development, as it is shown at **Listing 4.6**. However, all the recent versions of the major web browsers are compatible with it.

Listing 4.6 An example code, which demonstrates how one can check the support for IndexedDB API[25].

```
1 if (!('indexedDB' in window)) {
2   console.log('This browser doesn\'t support IndexedDB');
3   return;
4 }
```

Nevertheless, one of the most significant problems with IndexedDB is using it in the development. It has an asynchronous API, which is using events. That makes developers produce a complex code, which is hard to maintain. Therefore, in the design of WebCure, we are going to use

a wrapper over the IndexedDB API – IndexedDB Promised[26, 27]. It is a tiny library, written by Jake Archibald from Google, which makes the use of JavaScript promises and simplifies the development process with IndexedDB, while keeping its functionality.

4.4 Persistent Storage

Both Cache API and IndexedDB database mentioned in **Section 4.2** and **Section 4.3** are taking the place at the local machine. However, when the local machine is running out of storage space, user agents will clear it automatically on an LRU policy⁸[28]. This is not something that suits an application, which promises to provide offline experience. And in the worst case scenario, if the data was not synced with the server, it is going to be lost. The solution for this problem is to use a Persistent Storage API[29, 30], which guarantees that cached data is not going to be cleared if the browser comes under pressure.

4.5 Background Sync

The reality is that it is not always possible to be online all the time, even if someone wanted to. Sometimes, there is no network connection at all, or it could be that abysmal that it could be hard to do anything under such conditions. Therefore, the scenario when someone could do his or her work offline and then, when the connection is re-established again, this work will go online, is useful. Nowadays, thanks to *Background Sync*[31] from Google it is possible to do in web applications.

Listing 4.7 An example code, which demonstrates how to register a sync (*myFirstSync* here) event for the service worker[31].

```
1 // Register your service worker:
2 navigator.serviceWorker.register('/sw.js');

4 // Then later, request a one-off sync:
5 navigator.serviceWorker.ready.then(function(swRegistration) {
```

⁸In the systems with LRU or least recently used caching policy, the least frequently used data will be cleared first.

```

6   return swRegistration.sync.register('myFirstSync');
7 });

```

For this feature to work, the web application has to use a service worker. First of all, it should be registered and, afterwards, a unique tag should be registered as well, which is going to be responsible for the background call of the method. Let us show an example: at **Listing 4.7**, we can see how a synchronization tag *myFirstSync* is registered.

Listing 4.8 An example code, which demonstrates that a function *doSomeStuff* called, when the *sync* event happened[31].

```

1 self.addEventListener('sync', function(event) {
2   if (event.tag == 'myFirstSync') {
3     event.waitUntil(doSomeStuff());
4   }
5 });

```


Afterwards, when a page controlled by the service worker is going back online (which is when the user agent has established a network connection[32]), a *sync* event is triggered. There, it is possible to perform a distinct action, depending on the registered earlier tag. For instance, at **Listing 4.8** we can see that a function *doSomeStuff* is called once the connection is back, after a *sync* event occurred. This function has to return a promise, which could be successful or not. In the latter case, another sync will be scheduled to try when there is a connectivity[31].

Background synchronization is an advantageous feature, which can be used in different scenarios. It could be, for example, sending the e-mails or any other type of messages, after having failed to send it when the connection was poor. However, in our case, in WebCure we are going to use it for sending the operations conducted offline back to the AntidoteDB server for further synchronization. Moreover, if there was a use case with a requirement to get the most recent data from the server after the re-connection to the network, it is possible to implement it using the background synchronization feature.

5 Implementation

In this chapter, we are going to explain the major problems that we encountered during the development phase.

Addressing the client, in the **Chapter 4** we already named the key frameworks and techniques that we used in the implementation. Concerning the server, which is going to provide to the client the RESTful⁹ interface, we will cover the use cases.



Source: https://cdn-images-1.medium.com/max/660/1*EbBD6IXvf3o-YegUvRB_IA.jpeg

Figure 5.1: A look at the server-client architecture with a RESTful interface.

To demonstrate the functionality of the WebCure, we covered three different types of CRDTs used in AntidoteDB: a counter, a set and a multi-value register. We managed to support each of these data types on the server and the client as well.

⁹REST or Representational State Transfer is a set of design principles, which make the network communication more flexible and scalable[33].

5.1 System's main components

Let us have a look at the main components of the system:

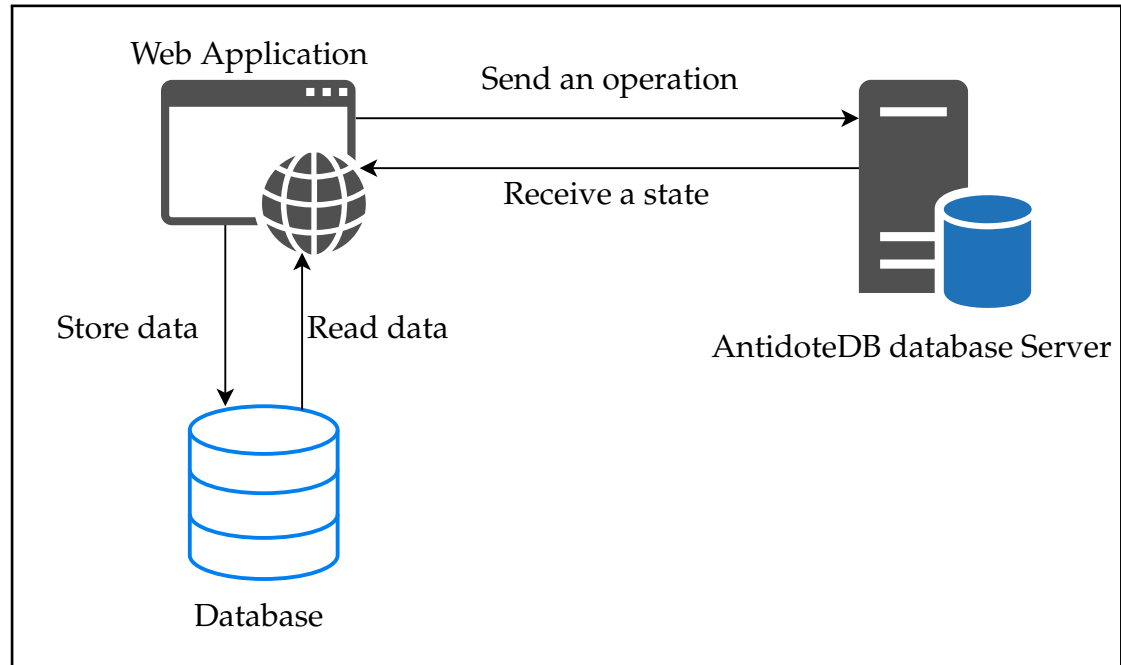


Figure 5.2: A top-level view of the system's design.

5.1.1 Web Application

First of all, in order for the user to be able to communicate with an AntidoteDB server, we are going to have a running web application that serves as a client. It runs in the web-browser, supports various commands from the user and sits on top of the local database layer.

For each of supported CRDTs, these are available commands:

- *read(key)* – an asynchronous function that pulls database changes concerning the *key* that the user passed.
 - *key*: a key to be read from the server's database;
 - *timestamp*: an optional parameter, which let us to set the timestamp at what the data should be read from;

- *update(key, op)* – an asynchronous function that processes user-made updates.
 - *key*: a key, which is going to be updated;
 - *op*: an operation, which is going to be performed on the key (it could differ depending on the data type it is running on);

When a user tries to update the data on a client-side, there are two possibilities:

- The request to the server is successful: in this case, the request would end up changing the data on the server's side, while the client will just update its own cache once it gets a response from the server.
- The request to the server is not successful: here, it is going to be a bit more interesting. We will have to wait for an error message and store the data in a temporary database for the updates that are not sent to the server's side just yet. Afterwards, I think it would be logical to have a timer, which will check the connection with the server. Once it is back, every transaction again is going to be sent to the server. After all the data is sent, this temporary database can be cleaned.

Several problems could arise, however, which we are going to discuss later.

5.1.2 Server

This part is an important part of the WebCure, as it provides the interface for the client and operates with the database layer as well. It supports the following scenarios for all the above-mentioned CRDTs:

- receiving an operation or an array of operations performed on a CRDT-object, according to the key, and applying them on the server;
- sending back to the client the state of requested CRDT-object / objects according to their state on the server;

5.1.3 A database layer

This layer should consist of the AntidoteDB database, which is going to store the actual states of CRDT-objects. When a user performs *read* by *key* operation from the cache, the following actions are taking place:

5 Implementation

- Firstly, the state of the object O is going to be found by *key* in the database
- Next, operations o performed on the object O are selected
- Afterwards, selected operations o applied on the object O .
- Finally, the object from the previous step is returned back as a response to the application.




Figure 5.3: 3-Tier architecture.

5.2 Optimization

6 Evaluation

Evaluation

This chapter is needed, when results from the previous chapters need to be systematically discussed. This is the case, when an implementation needs to be assessed, or the results of a case study or an experiment need to be interpreted.

- add the comparison between your approach and the other, in related work.
- add the test coverage statistics

7 Related Work

In this chapter, we are going to mention some of the researches, which influenced our work.

7.1 SwiftCloud: a transactional system that brings geo-replication to the client

Geo-replication of data into several data centers (DC) across the world is used in cloud platforms in order to improve availability and latency[34]. This goal could be achieved even to a greater extent by storing some part of the data or even by replicating all of it at client machines. Thus, caching could be useful in terms of increasing availability of systems.

One of such systems that integrates client- and server-side storage is called SwiftCloud, where the idea is to cache a subset of the objects from the DCs, and if the appropriate objects are in cache, then responsiveness is improved and the operation without an internet connection is possible[35]. The authors of the SwiftCloud state that it improves latency and throughput if it is compared to general geo-replication techniques. That is possible due to availability during faults thanks to automatic switch of DCs when the current one does not respond. Apart from that, the system also maintains consistency guarantees [7].

7.2 Legion: a framework, which enriches web applications

A framework named Legion shows another interesting approach how to address availability and scalability issues by using cache at the client-side. The idea is to avoid a concept of a centralized infrastructure for mediating user interactions, as it causes unnecessarily high latency and hinders fault-tolerance and scalability[36]. As an alternative, authors of Legion suggest client web applications to securely replicate data from servers and afterwards synchronize the data among the clients. This change makes the

system less dependent on the server and, moreover, it reduces the latency of interactions among the clients. The guarantee of convergence between all of the replicas is possible due to CRDTs, introduced in [Chapter 2](#).

7.3 Developing web and mobile applications with offline usage experience

Nowadays, applications with offline experience are becoming extremely popular. In this thesis, in [Chapter 4](#) we presented an approach of developing web applications with a help of service workers and background synchronization (they also go by name of Progressive Web Applications [PWA]). It is a universal approach to create a cross-platform application that would work in web and on mobile devices. However, there are also other ways how the offline experience could be achieved. Sometimes, the process could become easier if specific frameworks are used. In the following subsections, we are going to describe tools and techniques that are useful in this context.

Polymer App Toolbox

Polymer is a JavaScript library from Google, which helps to build web applications with the use of Web Components. The former concept represents a set of web platform APIs, which allow to create custom HTML tags to use in web pages[37]. As web components are based on the latest web standards and it eases the process of development. Polymer App Toolbox, in its turn, provides a collection of components to build PWAs with Polymer. However, the support of offline-experience is possible yet again due to Service Workers[38], which repeats the solution used in this thesis. Currently, among the users of Polymer, apart from Google, there are such giants as Electronic Arts, IBM and Coca-Cola[39].

HTML5 Specification

The specification of HTML5 contains some features that offer a possibility to build web applications that work offline. The solution to address this problem is to use SQL-based database API in order to be able to store data locally, and to use an offline application HTTP cache to make sure about the availability of the application when there is no internet connection.

The latter makes possible the following advantages: offline browsing, flexibility and fast load of resources from the hard drive[40]. However, from 2015 the application cache is considered to be deprecated and is recommended to be avoided[41] in favour of Service Workers. Apart from that, HTML5 also lets the developer to check in which stage the user's navigator is, by simply checking the attribute *navigator.onLine*[42]. Offline and online events make changes to this attribute, which reflects the status of current connection. The advantages of HTML5 approach is the reusability of code, shorter development and testing times and universal support of mobile platforms. However, if compared to the mobile applications written in native languages, HTML5 web applications could be inferior to them.

Hoodie

Hoodie is a framework, which eases the process of developing web and iOS applications. We will not cover all the features it offers and only stop on it providing offline experience for applications that are developed using Hoodie. The documentation states that framework is offline-first[43], which means that all the data is firstly stored locally and could be accessed without the internet connection. This is possible due to PouchDB database, which performs this work in the background[44]. We will explore the process of how PouchDB works below.

localStorage

localStorage is a JavaScript library, which provides the possibility to improve the offline-experience of web applications in terms of storing data on the client-side. It uses IndexedDB, localStorage or WebSQL with a simple API. localStorage sits on top of datastore layer and provides a range of methods to control the data. One of the useful benefits of it is that the data is not required to be explicitly converted into JSON format, as localStorage does that automatically[45].

PouchDB and CouchDB

PouchDB represents an open-source JavaScript database, which is designed to build applications, which work well online and offline. To put it in a nutshell, PouchDB enables web applications to store the data locally,

and, apart from that, eases the process of synchronization with CouchDB-compatible servers. CouchDB, in its turn, is a database that is supported by a replication approach, which allows to synchronize two or more servers, based on CouchDB. As there is a replication approach, it has its own way of dealing with conflicts, which we explain next by following the description of the protocol offered by Lehnardt [46]. For every item stored in CouchDB, the database will add two extra properties: *_id* and *_rev*, which can be seen at **Listing 7.1** that shows the result of getting an element named *document* previously stored in the database.

Listing 7.1 A typical result of retrieving the item *document* stored in CouchDB.

```
1 {"_id":"document", "_rev":"1-23202479633c2b380f79507a776743d5", "a":1}
```

As you can see, the *_id* represents the name of the item, which is a custom name set by the user, while *_rev* represents the hash value, associated with the content. Afterwards, whenever we want to change the item *document*, the same *_id* and *_rev* should be used. For example, to change the value of *document* by adding *b* into it, a simple PUT HTTP-request should be sent, as **Listing 7.2** shows.

Listing 7.2 Updating the value of item *document* by adding *b* into it.

```
1 PUT /database/document
2 {"_id":"document", "_rev":"1-23202479633c2b380f79507a776743d5", "a":1, "b":2}
```

The next retrieval of the *document* will get us the result shown at **Listing 7.3**.

Listing 7.3 The result of requesting the updated version of *document*

7.3 Developing web and mobile applications with offline usage experience

```
1 GET /database/document
2 { "_id": "document", "_rev": "2-c5242a69558bf0c24dda59b585d1a52b"
  , "a": 1, "b": 2 }
```

As you can see, the `_rev` property got updated. In case the wrong revision is sent to update the document, the database will respond with an error.

However, as you might have guessed, it is much more interesting, when you have more than one server. Let us assume that now we have two CouchDB servers. Imagine we try to update the *document* once again with the values shown at **Listing 7.4**.

Listing 7.4 Updating the value of item *document* by adding *d* into it.

```
1 PUT /database/document
2 { "_id": "document", "_rev": "2-c5242a69558bf0c24dda59b585d1a52b"
  , "a": 1, "b": 2, "d": 4 }
```

For example, let us say that it gets written to the first CouchDB server, which gives the response shown at **Listing 7.5**.

Listing 7.5 The result of requesting the *document* from CouchDB-1.

```
1 GET /database/document
2 { "_id": "document", "_rev": "3-2235fd4815b81b2da1b84159aba4006e"
  , "a": 1, "b": 2, "d": 4 }
```

But the second CouchDB server still has the old data, as you can see at **Listing 7.6**.

Listing 7.6 The result of requesting the *document* from CouchDB-2.

7 Related Work

```
1 GET /database/document
2 {"_id":"document","_rev":"2-c5242a69558bf0c24dda59b585d1a52b"
  , "a":1, "b":2}
```

Ideally, the replication happens fastly and both databases will synchornize and reach the same state. However, sometimes it might take a while. And if a client tries to update the document yet another time, there is a possibility that the update will go to the second server, which will reject the update, as *_rev* does not match any more.

Listing 7.7 Updating the value of item *document* by adding element *e* and removing previously added element *d*

```
1 PUT /database/document
2 {"_id":"document","_rev":"3-2235fd4815b81b2da1b84159aba4006e"
  , "a":1, "b":2, "e":5}
```

As we mentioned ealier, the response of the second CouchDB server for the operation shown at **Listing 7.7** will look like the one at **Listing 7.8**.

Listing 7.8 The demonstration of a conflict situation happening, when the *_rev* of sent operation and the one at the server do not match.

```
1 {"error":"conflict","reason":"Document update conflict."}
```

However, there is another option to perform this update. We might have a strategy of getting the latest *_rev* first, which was "2-c5242a69558bf0c24dda59b585d1a52b" at the second CouchDB server, and only then applying the update. So, the operation will look as at **Listing 7.9**.

Listing 7.9 Updating the value of item *document* by adding element *e* and removing previously added element *d* after receiving the new *_rev* from the second CouchDB server.

```
1 PUT /database/document
2 { "_id": "document", "_rev": "2-c5242a69558bf0c24dda59b585d1a52b"
  , "a": 1, "b": 2, "e": 5 }
```

In case this PUT request goes to the second CouchDB server, then this operation will be successful and now both servers will have different states, which creates a conflict situation. However, it is still an undesirable situation. Thus, there are the limitations that should be given a thought in the development process, which indeed make the process of creating the product based on CouchDB more complicated:

- Making a change, do not request multiple GETs and POSTs
- Do not update the *_rev* locally in the client without getting new data from the server before that

Realm Mobile

This is a framework, which makes an integration of a client-side database for iOS and Android with a server-side, which offers the following features: real-time synchronization, conflict resolution and event handling. This framework eases the process of developing applications with offline experience. The concept we are interested here is conflict-handling. In AntidoteDB, for this purpose, CRDTs are used. Realm Mobile maintains a good user experience in offline mode due to the rules they described for conflict resolution. As they state, “at a very high level the rules are as follows:

- **Deletes always win.** If one side deletes an object it will always stay deleted, even if the other side has made changes to it later on.
- **Last update wins.** If two sides update the same property, the value will end up as the last updated.
- **Inserts in lists are ordered by time.** If two items are inserted at the same position, the item that was inserted first will end up before the other item. This means that if both sides append items to the end of a list they will end up in order of insertion time.”[47]

7 *Related Work*

The authors of the framework state that “strong eventual consistency” guarantees are reached[47] with the above mentioned approach. Though, such way of conflict handling is not as flexible as CRDTs and has to be taken into account by the programmer at the stage of the development.

8 Conclusion

Conclusion

This is always the first chapter of the thesis. The chapter should be short (up to 5 pages). The chapter should feature sections as follows (where applicable):

- o Summary (Summarize this work in an insightful manner, assuming that the reader has seen the rest.)
- o Limitations or threats to validity (Point out the limitations of this work. In the case of empirical research, discuss threats to validity in a systematic manner.)
- o Future work (Provide insightful advice on where this research should be taken next.)

8.1 Summary

8.2 Future Work

- Automatic synchronization of updates like without pressing any buttons;
- The thing that was mentioned in the design chapter about the state-based approach;

Idempotence of updates

- wrote about this part in the notes to the architecture pdf

In case a client is sending an update to the server and does not receive an acknowledgement, what could have happened?

There are possibilities:

- The update was applied on the server, but the connection failed when the acknowledgement was about to be sent;
- The update was not applied on the server and the client did not receive an acknowledgement;

However, the client does not know which of these situations happened. Therefore, we have to find a general solution for such kind of behaviour.

One of the solutions could be the following: it does not matter, whether the update was applied or not. A client will just send the update again,

8 Conclusion

till it does not receive an acknowledgement, regardless of what happens on the server's side.

The other one is: the client sends an update to the server, which should send a message back that it received an update. Afterwards, a client removes this update from the temporary database and sends back a message to the server that it is possible to apply the update.

However, the implementation of the solution for this problem is beyond the scope of this Master thesis. Therefore, these thoughts are going to be included in a section, where future improvements will be discussed.

List of Figures

2.1	An example of how a causally consistent behaviour (a) and the one that is not (b) could work in a social network. . . .	5
2.2	Operation-based approach[12]. «S» stands for source replicas and «D» for downstream replicas.	7
2.3	State-based approach[12]. «S» stands for source replicas and «M» for merging stages.	8
3.1	An overview of the communication protocol.	14
3.2	An overview of notations used in the following chapters for the protocol explanation.	14
3.3	The way a client receives a key-value update from the server.	15
3.4	The way a client sends a write operation to the server. . . .	16
3.5	Graphical representation over a possible communication between two clients and a server.	17
3.6	Successful request of state from the local database – Sequence diagram.	18
3.7	Successful storing of an operation in the local database – Sequence diagram.	19
4.1	An overview of the service worker being able to provide the application experience of both online and offline modes.	22
5.1	A look at the server-client architecture with a RESTful interface.	31
5.2	A top-level view of the system's design.	32
5.3	3-Tier architecture.	34

List of Tables

3.1	Functional requirements.	12
3.2	Non-functional requirements.	12

List of Code

Listings

3.1	A pseudocode for requesting the data on a client's side. . .	15
4.1	An example code, which demonstrates how to register a service worker[19].	23
4.2	An example code, which demonstrates how to set a custom scope when registering a service worker[19].	24
4.3	An example code, which demonstrates a listener for the <i>install</i> service worker's event[19].	25
4.4	An example code, which demonstrates a listener for the <i>activation</i> service worker's event[19].	25
4.5	An example code, which demonstrates how one can create cache storage called <i>my-cache</i> [23].	26
4.6	An example code, which demonstrates how one can check the support for IndexedDB API[25].	27
4.7	An example code, which demonstrates how to register a sync (<i>myFirstSync</i> here) event for the service worker[31]. .	28
4.8	An example code, which demonstrates that a function <i>doSomething</i> called, when the <i>sync</i> event happened[31].	29
7.1	A typical result of retrieving the item <i>document</i> stored in CouchDB.	40
7.2	Updating the value of item <i>document</i> by adding <i>b</i> into it. .	40
7.3	The result of requesting the updated version of <i>document</i> .	40
7.4	Updating the value of item <i>document</i> by adding <i>d</i> into it. .	41
7.5	The result of requesting the <i>document</i> from CouchDB-1. . .	41
7.6	The result of requesting the <i>document</i> from CouchDB-2. . .	41
7.7	Updating the value of item <i>document</i> by adding element <i>e</i> and removing previously added element <i>d</i>	42
7.8	The demonstration of a conflict situation happening, when the <i>_rev</i> of sent operation and the one at the server do not match.	42

Listings

- 7.9 Updating the value of item *document* by adding element *e* and removing previously added element *d* after receiving the new *_rev* from the second CouchDB server. 42

Bibliography

- [1] M. Tamer zsu and Patrick Valduriez. *Principles of Distributed Database Systems*. Springer Publishing Company, Incorporated, 3rd edition, 2011. ISBN 1441988335, 9781441988331.
- [2] Hagit Attiya and Jennifer L. Welch. Sequential consistency versus linearizability. *ACM Trans. Comput. Syst.*, 12(2):91–122, May 1994. ISSN 0734-2071. doi: 10.1145/176575.176576. URL <http://doi.acm.org/10.1145/176575.176576>.
- [3] Irene Y. Zhang. Consistency should be more consistent! <https://irenezhang.net/blog/2015/02/01/consistency.html>, 2015 (accessed October 23, 2018).
- [4] Marc Shapiro. Strong eventual consistency and conflict-free replicated data types. <https://www.microsoft.com/en-us/research/video/strong-eventual-consistency-and-conflict-free-replicated-data-types/>, 2011 (accessed October 23, 2018).
- [5] S. Gilbert and N. Lynch. Perspectives on the cap theorem. *Computer*, 45(2):30–36, Feb 2012. ISSN 0018-9162. doi: 10.1109/MC.2011.389.
- [6] Santiago J. Castiñeira and Annette Bieniusa. Collaborative offline web applications using conflict-free replicated data types. In *Proceedings of the First Workshop on Principles and Practice of Consistency for Distributed Data, PaPoC '15*, pages 5:1–5:4, New York, NY, USA, 2015. ACM. ISBN 978-1-4503-3537-9. doi: 10.1145/2745947.2745952. URL <http://doi.acm.org/10.1145/2745947.2745952>.
- [7] Marek Zawirski, Nuno Preguiça, Sérgio Duarte, Annette Bieniusa, Valter Balegas, and Marc Shapiro. Write fast, read in the past: Causal consistency for client-side applications. In *Proceedings of the 16th Annual Middleware Conference, Middleware '15*, 2015. ISBN 978-1-4503-3618-5. doi: 10.1145/2814576.2814733. URL <http://doi.acm.org/10.1145/2814576.2814733>.

Bibliography

- [8] Antidotedb. <https://www.antidotedb.eu/>, 2018 (accessed September 25, 2018).
- [9] Annette Bieniusa Deepthi Devaki Akkoorath. Antidote: the highly-available geo-replicated database with strongest guarantees. *SyncFree Technology White Paper*, 2016.
- [10] Butler W. Lampson. Atomic transactions. In *Distributed Systems - Architecture and Implementation, An Advanced Course*, 1981. ISBN 3-540-10571-9. URL <http://dl.acm.org/citation.cfm?id=647434.726386>.
- [11] Deepthi Devaki Akkoorath, Alejandro Z. Tomsic, Manuel Bravo, Zhongmiao Li, Tyler Crain, Annette Bieniusa, Nuno Preguica, and Marc Shapiro. Cure: Strong Semantics Meets High Availability and Low Latency. *Proceedings - International Conference on Distributed Computing Systems*, 2016. ISSN 1063-6927. doi: 10.1109/ICDCS.2016.98.
- [12] Marc Shapiro, Nuno Pregui, Carlos Baquero, and Marek Zawirski. Replicated Data Types: A comprehensive study of Convergent and Commutative Replicated Data Types. 2011. URL <http://hal.inria.fr/inria-00555588>.
- [13] Nuno Preguiça, Carlos Baquero, and Marc Shapiro. *Conflict-Free Replicated Data Types CRDTs*. Springer International Publishing, 2018. ISBN 978-3-319-63962-8. doi: 10.1007/978-3-319-63962-8_185-1. URL https://doi.org/10.1007/978-3-319-63962-8_185-1.
- [14] Datatypes in antidote. <https://antidotedb.gitbook.io/documentation/api/datatypes>, 2018 (accessed November 26, 2018).
- [15] Black box model. <https://www.investopedia.com/terms/b/blackbox.asp>, 2018 (accessed November 29, 2018).
- [16] Caching files with service worker. <https://developers.google.com/web/ilt/pwa/caching-files-with-service-worker>, 2018 (accessed November 28, 2018).
- [17] Matt Gaunt. Service workers: an introduction. <https://developers.google.com/web/fundamentals/primers/service-workers/>, 2018 (accessed September 14, 2018).
- [18] Web workers api. https://developer.mozilla.org/en-US/docs/Web/API/Web_Workers_API, 2018 (accessed November 15, 2018).

- [19] Introduction to service worker. <https://developers.google.com/web/ilt/pwa/introduction-to-service-worker>, 2018 (accessed November 18, 2018).
- [20] Yvo Desmedt. *Man-in-the-Middle Attack*, pages 368–368. Springer US, Boston, MA, 2005. ISBN 978-0-387-23483-0. doi: 10.1007/0-387-23483-7_241. URL https://doi.org/10.1007/0-387-23483-7_241.
- [21] The navigator object. https://www.w3schools.com/jsref/obj_navigator.asp, 2018 (accessed November 28, 2018).
- [22] Service worker api. https://developer.mozilla.org/en-US/docs/Web/API/Service_Worker_API, 2018 (accessed November 28, 2018).
- [23] Mat Scales. Using the cache api. <https://developers.google.com/web/fundamentals/instant-and-offline/web-storage/cache-api>, 2018 (accessed November 28, 2018).
- [24] Promise. https://developer.mozilla.org/en-US/docs/Web/JavaScript/Reference/Global_Objects/Promise, 2018 (accessed November 23, 2018).
- [25] Working with indexeddb. <https://developers.google.com/web/ilt/pwa/working-with-indexeddb>, 2018 (accessed November 23, 2018).
- [26] Jake Archibald. Indexeddb promised. <https://github.com/jakearchibald/idb>, 2018 (accessed November 23, 2018).
- [27] Lab: Indexeddb. <https://developers.google.com/web/ilt/pwa/lab-indexeddb>, 2018 (accessed November 29, 2018).
- [28] Browser storage limits and eviction criteria. https://developer.mozilla.org/en-US/docs/Web/API/IndexedDB_API/Browser_storage_limits_and_eviction_criteria#LRU_policy, 2018 (accessed November 29, 2018).
- [29] Chris Wilson. Persistent storage. <https://developers.google.com/web/updates/2016/06/persistent-storage>, 2018 (accessed November 28, 2018).
- [30] Dean Hume. Persistent storage api: Building for the offline web. <https://deanhume.com/persistent-storage-api-building-for-the-offline-web/>, 2017 (accessed November 28, 2018).

Bibliography

- [31] Jake Archibald. Introducing background sync. <https://developers.google.com/web/updates/2015/12/background-sync>, 2018 (accessed November 24, 2018).
- [32] Marijn Kruisselbrink Josh Karlin. Web background synchronization. <https://wicg.github.io/BackgroundSync/spec/#online>, 2018 (accessed November 28, 2018).
- [33] Lauren Long. What restful actually means. <https://codewords.recurse.com/issues/five/what-restful-actually-means>, 2018 (accessed November 27, 2018).
- [34] Yair Sovran, Russell Power, Marcos K. Aguilera, and Jinyang Li. Transactional storage for geo-replicated systems. In *Proceedings of the Twenty-Third ACM Symposium on Operating Systems Principles, SOSP '11*, 2011. ISBN 978-1-4503-0977-6. doi: 10.1145/2043556.2043592. URL <http://doi.acm.org/10.1145/2043556.2043592>.
- [35] Nuno Preguiça, Marek Zawirski, Annette Bieniusa, Sérgio Duarte, Valter Balegas, Carlos Baquero, and Marc Shapiro. Swiftcloud: Fault-tolerant geo-replication integrated all the way to the client machine. In *Proceedings of the 2014 IEEE 33rd International Symposium on Reliable Distributed Systems Workshops, SRDSW '14*, 2014. ISBN 978-1-4799-7361-3. doi: 10.1109/SRDSW.2014.33. URL <https://doi.org/10.1109/SRDSW.2014.33>.
- [36] Albert van der Linde, Pedro Fouto, João Leitão, Nuno Preguiça, Santiago Castiñeira, and Annette Bieniusa. Legion: Enriching Internet Services with Peer-to-Peer Interactions. *Proceedings of the 26th International Conference on World Wide Web*, 2017. doi: 10.1145/3038912.3052673.
- [37] What are web components? <https://www.webcomponents.org/introduction#what-are-web-components->, 2018 (accessed October 29, 2018).
- [38] Offline caching with service worker precache. <https://www.polymer-project.org/1.0/toolbox/service-worker>, 2016 (accessed October 29, 2018).
- [39] Who's using polymer? <https://github.com/Polymer/polymer/wiki/Who's-using-Polymer%3F>, 2018 (accessed October 29, 2018).

- [40] Eric Bidelman. A beginner's guide to using the application cache. <https://www.html5rocks.com/en/tutorials/appcache/beginner/>, 2018 (accessed October 30, 2018).
- [41] Using the application cache. https://developer.mozilla.org/en-US/docs/Web/HTML/Using_the_application_cache, 2018 (accessed October 30, 2018).
- [42] Anne van Kesteren. Offline web applications. <https://www.w3.org/TR/offline-webapps/>, 2008 (accessed October 29, 2018).
- [43] hood.ie intro. <http://hood.ie/intro/>, accessed November 4, 2018.
- [44] Jake Peyser. Intro to hoodie and react. <https://css-tricks.com/intro-hoodie-react/>, 2017 (accessed November 4, 2018).
- [45] Matt West. Using localforage for offline data storage. <https://blog.teamtreehouse.com/using-localforage-offline-data-storage>, 2013 (accessed November 4, 2018).
- [46] Jan Lehnardt. Understanding couchdb conflicts. <https://writing.jan.io/2013/12/19/understanding-couchdb-conflicts.html>, 2013 (accessed November 4, 2018).
- [47] Conflict resolution. <https://docs.realm.io/platform/self-hosted/customize/conflict-resolution>, 2018 (accessed October 30, 2018).